

TEXT MINING

Lecture 02

INTERMEDIATE R

KEUNGOU I KIM
awekim@handong.edu



R Basics

Package Installation and Loading

- `install.packages("package name")`

```
> install.packages('rpart')
```

```
trying URL 'https://cran.rstudio.com/bin/windows/contrib/4.1/rpart_4.1.16.zip'
```

```
Content type 'application/zip' length 983052 bytes (960 KB)
```

```
downloaded 960 KB
```

```
package 'rpart' successfully unpacked and MD5 sums checked
```

```
The downloaded binary packages are in
```

```
C:\Users\aweki\AppData\Local\Temp\Rtmpum62C9\downloaded_packages
```

- to load package

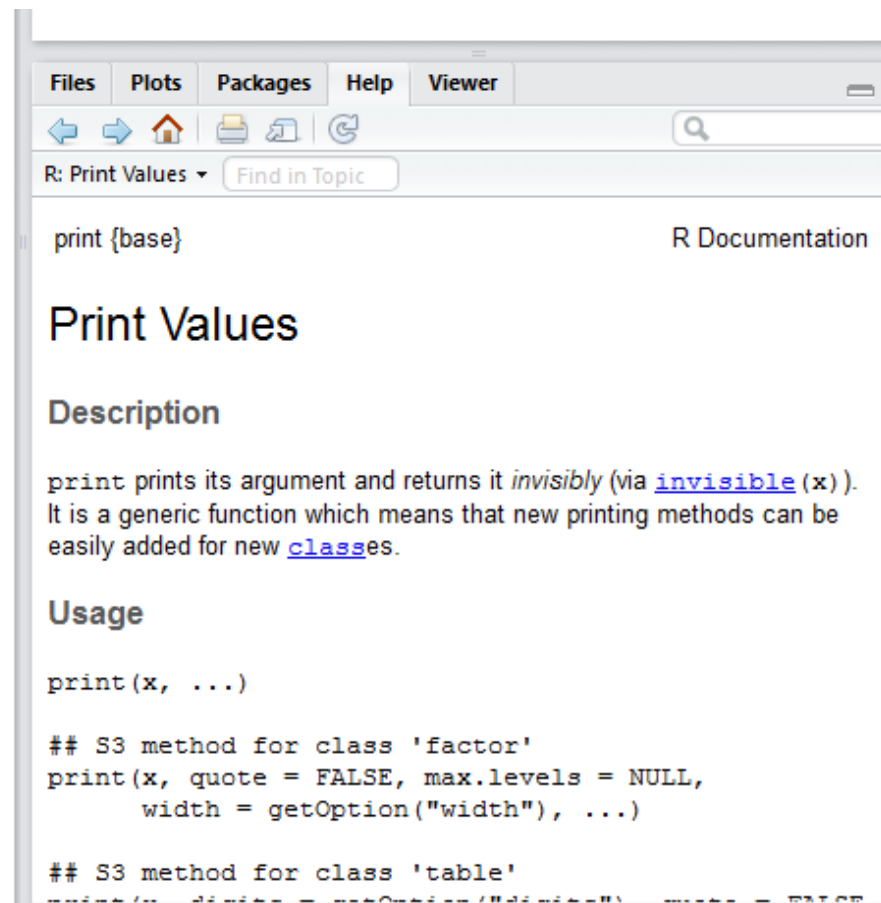
- `library("package name")`

- or `require("package name")`

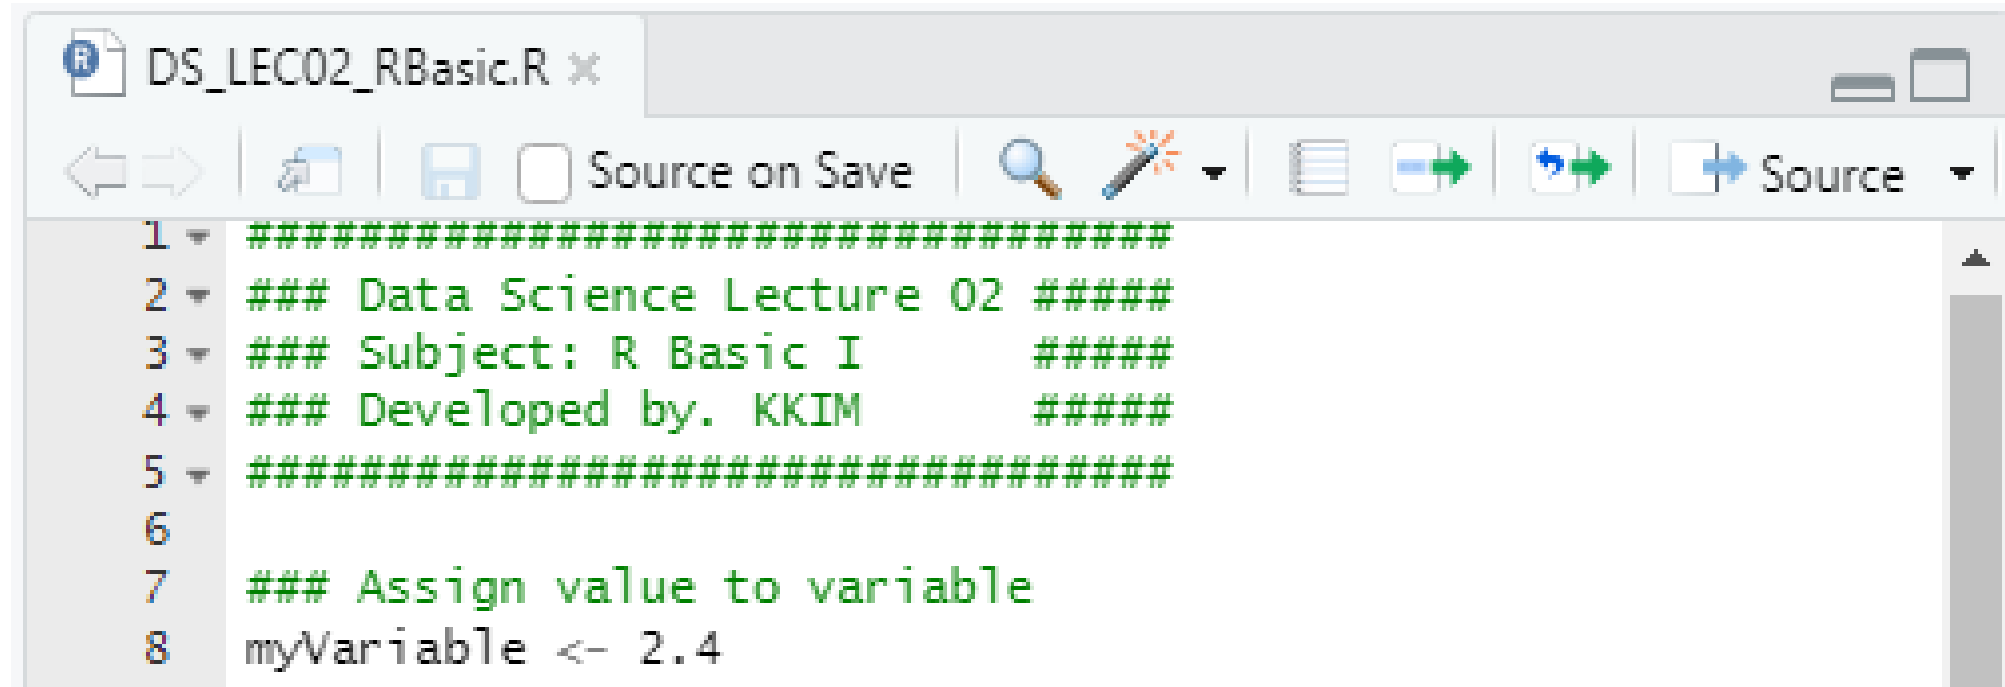
- Getting help
 - `help(command)` or `?command`
 - `example(command)` to see examples

'help.start()' for an HTML browser interface to help.
Type 'q()' to quit R.

```
> print("hello")  
[1] "hello"  
> a = 10  
> b = 20  
> a+b  
[1] 30  
> ?print  
> |
```



- Comments are not operated in R
 - It is used to describe the program author, date, and nature of the code.
 - It is used to help programmers understand the code
- # or ctrl + shift + c



The screenshot shows an R script editor window titled "DS_LEC02_RBASIC.R". The editor contains the following code:

```
1 #####  
2 ### Data Science Lecture 02 ###  
3 ### Subject: R Basic I #####  
4 ### Developed by. KKIM #####  
5 #####  
6  
7 ### Assign value to variable  
8 myVariable <- 2.4
```

- A vector is a **sequence of data elements**.
 - A set with multiple elements
 - “sequence” meaning that it has an order, or it has an index number
→ Indexing and Slicing can be used
- `c()` function
 - A generic function that combines its arguments
→ fundamental function for creating a “vector”
 - Used for combining elements

```
> c("Lee", "Yoon", "Shim", "Ahn", "Oh", "Huh")  
[1] "Lee" "Yoon" "Shim" "Ahn" "Oh" "Huh"  
> PresCand <- c("Lee", "Yoon", "Shim", "Ahn", "Oh", "Huh")  
> PresCand  
[1] "Lee" "Yoon" "Shim" "Ahn" "Oh" "Huh"
```

- Conventional data set with rows and columns
 - e.g. student datasets may contain name(character), age(integer), major(factor), GPA(numeric, real number)...
 - Vectors and metrics can have values of the same data type
- A DataFrame has the variables of a data set as columns and the observations as rows.
 - “List of Vectors”
 - DataFrames can have variables(vectors) of the same length (and possibly different types)
- data.frame() function

- A list in R allows you to gather a variety of objects under one name (that is, the name of the list) in an ordered way.
 - List objects can be matrices, vectors, data frames, even other lists, etc.
 - It is not even required that these objects are related to each other in any way.
 - Fewer restrictions than the DataFrame
- list() function

R Indexing & Slicing

- Data structure
 - 0D: Scala
 - 1D: Vector
 - 2D: Matrix, DataFrame
 - 3D: List
- Indexing
 - Way of accessing specific value
- Slicing
 - Way of accessing specific values

Indexing & Slicing

- 0D (scalar)

- `var0 <- 0.1`
- No need for indexing and slicing

```
> var0 <- 0.1  
> var0  
[1] 0.1
```

- 1D (vector)

- `var1 <- c(2, 4, 6, 8, 20)`

var1 contains 5 elements.

```
> var1 <- c(2, 4, 6, 8, 20)  
> var1  
[1] 2 4 6 8 20
```

- How can we access the 3rd element from var1?

```
> var1[3]  
[1] 6
```

`vec [IndexNumber]`

- How can we access the 1st – 3rd element from var1?

```
> var1[1:3]  
[1] 2 4 6
```

`vec [StartIndex : EndIndex]`

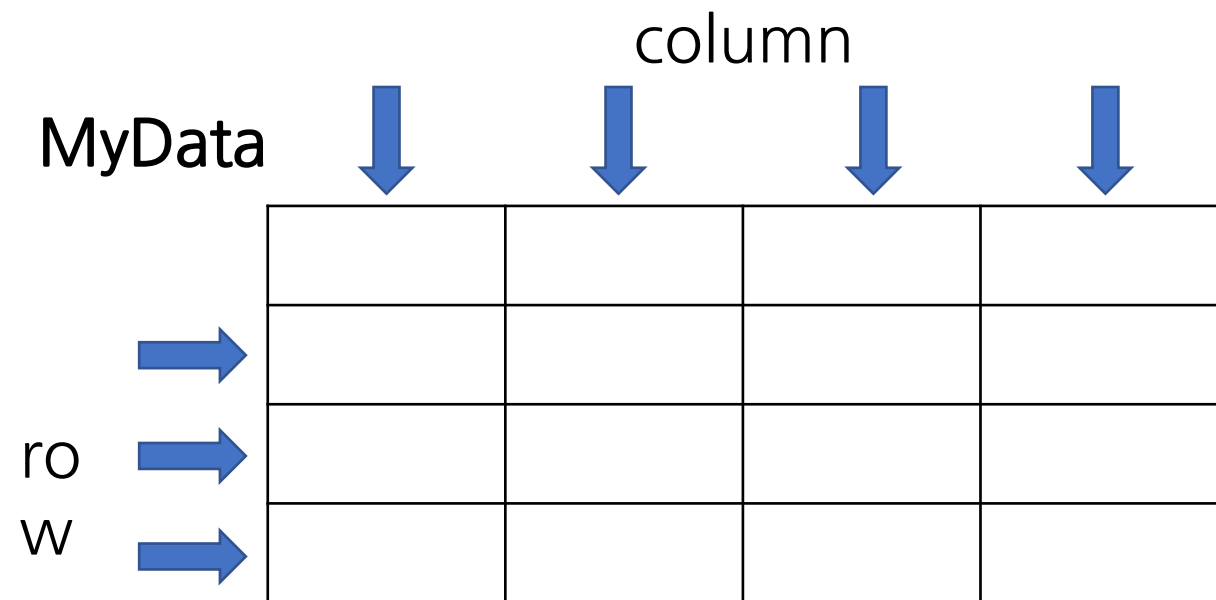
- How can we access the 1st & 3rd element from var1?

```
> var1[c(1,3)]  
[1] 2 6
```

`vec [c(IndexNumber1, IndexNumber2, ...)]`

Indexing & Slicing

- 2D (matrix, DataFrame)
 - Method 1) Numbers



DF [row-index, column-index]

*Either use
index number or slicing*

```
> data(women)
> women
   height weight
1      58    115
2      59    117
3      60    120
4      61    123
5      62    126
6      63    129
7      64    132
8      65    135
9      66    139
10     67    142
11     68    146
12     69    150
13     70    154
14     71    159
15     72    164
```

```
> women[1,1]
```

```
[1] 58
```

```
> women[1,]
```

```
   height weight
```

```
1      58    115
```

```
> women[,1]
```

```
[1] 58 59 60 61 62 63 64 65 66 67 68 69 70 71 72
```

```
> women[1:2,1]
```

```
[1] 58 59
```

```
> women[1:2,]
```

```
   height weight
```

```
1      58    115
```

```
2      59    117
```

Indexing & Slicing

- 2D (matrix, DataFrame)
 - Method 1) Names (+Number)

column

MyData

A 4x4 matrix diagram. Three blue arrows point downwards from the word 'column' to the top row of the matrix. Three blue arrows point to the right from the word 'row' to the first column of the matrix.

DF[row-name, column-name]

DF\$column-name[column-number]

DF\$column-name[column-slicing]

```
> data(women)
> women
      height weight
1         58    115
2         59    117
3         60    120
4         61    123
5         62    126
6         63    129
7         64    132
8         65    135
9         66    139
10        67    142
11        68    146
12        69    150
13        70    154
14        71    159
15        72    164
```

```
> women[1,c("height")]
```

```
[1] 58
```

```
> women[1:2,c("height","weight")]
```

```
      height weight
1         58    115
2         59    117
```

```
> women$height[1]
```

```
[1] 58
```

```
> women$height[1:2]
```

```
[1] 58 59
```

R Coding Overview

Code Writing

- When writing or reading programming codes, remember ...
 - Left → Right
 - Top → Down
 - In → Out (when [] or () is used)

```
DS_LEC02_RBasicI.R x
Source on Save
Run
Source

1 #####
2 ### Data Science Lecture 02 #####
3 ### Subject: R Basic I #####
4 ### Developed by. KKIM #####
5 #####
6
7 ### Assign value to variable
8 myVariable <- 2.4
9
10 ### Data Type
11 typeof(TRUE)
12 typeof("Hello")
13 typeof(3.14)
14 typeof(1L)
15
16 ### Operators
17 a <- 10.5
18 b <- 20
19 c <- 4
20
21 a + b ## addition
22 a - c ## subtraction
23 a * c ## multiplication
24 b / c ## division
25 a %% c ## remainder
26 a > b ## inequality
27 a*2 == b ## equality
28 1/(a < b) ## negation
```

head(AA[order(AA \$Var, decreasing = TRUE),])

*We can “understand” the code,
but it will take some time...*

Pipe Operator in R

- Also known as “chain operator”



- Can be used with the ‘dplyr’ package
- Using the pipe operator, we can write and read codes from left to write. No more inside to outside

```
> head(women)
  height weight
1     58   115
2     59   117
3     60   120
4     61   123
5     62   126
6     63   129
```

```
> women %>% head
  height weight
1     58   115
2     59   117
3     60   120
4     61   123
5     62   126
6     63   129
```

*Women is used as
an input for head()
function*

- Pipe operators
 - (dplyr package) %>%
 - (magrittr package) %<>%: Similar to %>%, but it “updates” the input

%>%

```
> women %>% head
```

	height	weight
--	--------	--------

1	58	115
2	59	117
3	60	120
4	61	123
5	62	126
6	63	129

```
> women
```

	height	weight
--	--------	--------

1	58	115
2	59	117
3	60	120
4	61	123
5	62	126
6	63	129
7	64	132
8	65	135
9	66	139
10	67	142
11	68	146
12	69	150
13	70	154
14	71	159
15	72	164

%<>%

```
> women %<>% head
```

```
> women
```

	height	weight
--	--------	--------

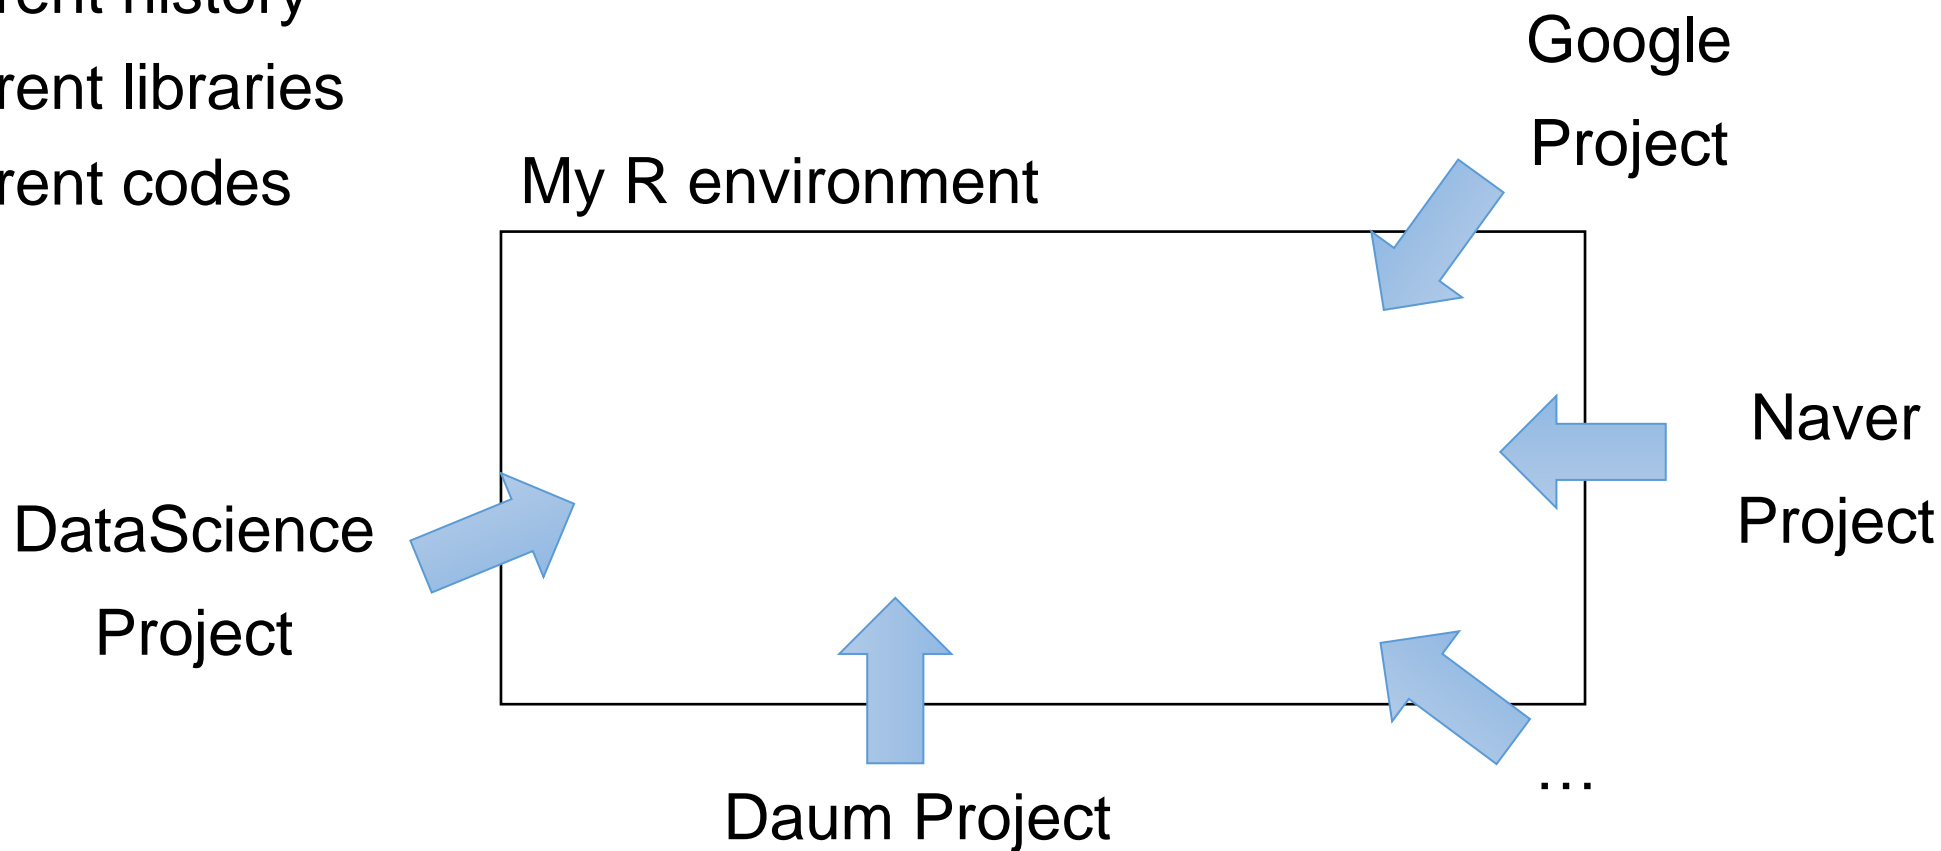
1	58	115
2	59	117
3	60	120
4	61	123
5	62	126
6	63	129

Advantages of Using Pipe Operator


- Readability
 - Read from left to right
- Continuity
 - Write from left to right
- In Python, the pipe operator is not needed.
 - Object-oriented programming → class and instance
 - Also known as method chaining or flow programming

R Project & Working Environment

- To conduct various data analysis projects, it is necessary to "well" manage each project
 - Different data
 - Different history
 - Different libraries
 - Different codes

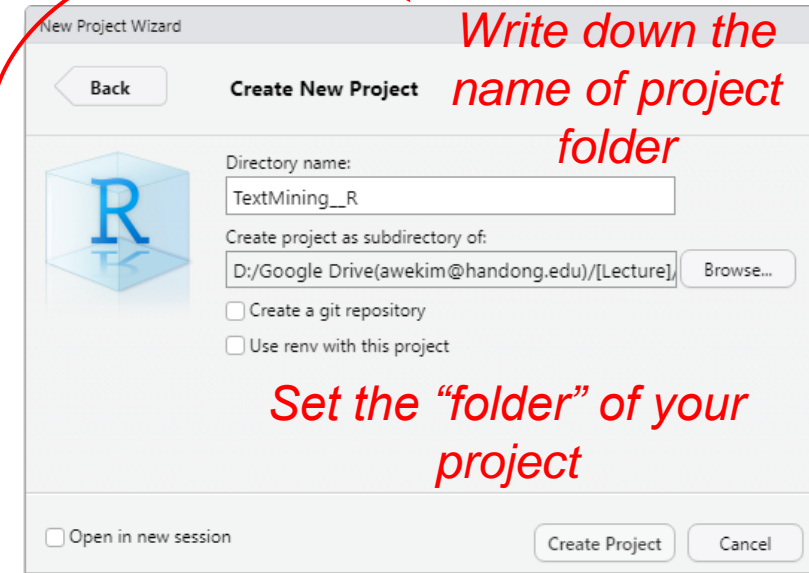
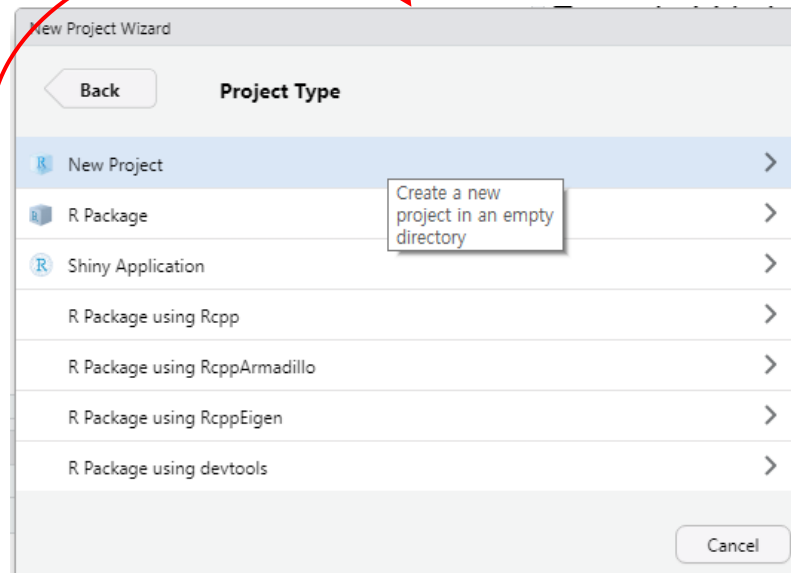
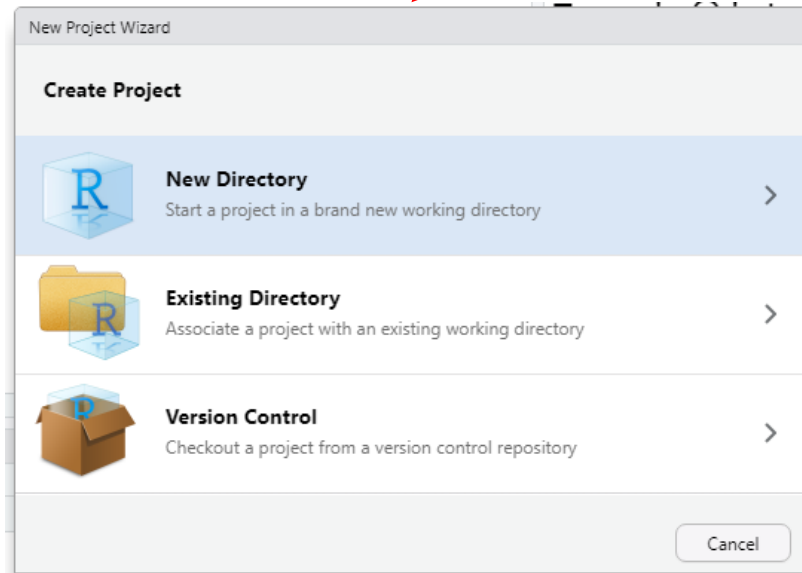


- R project



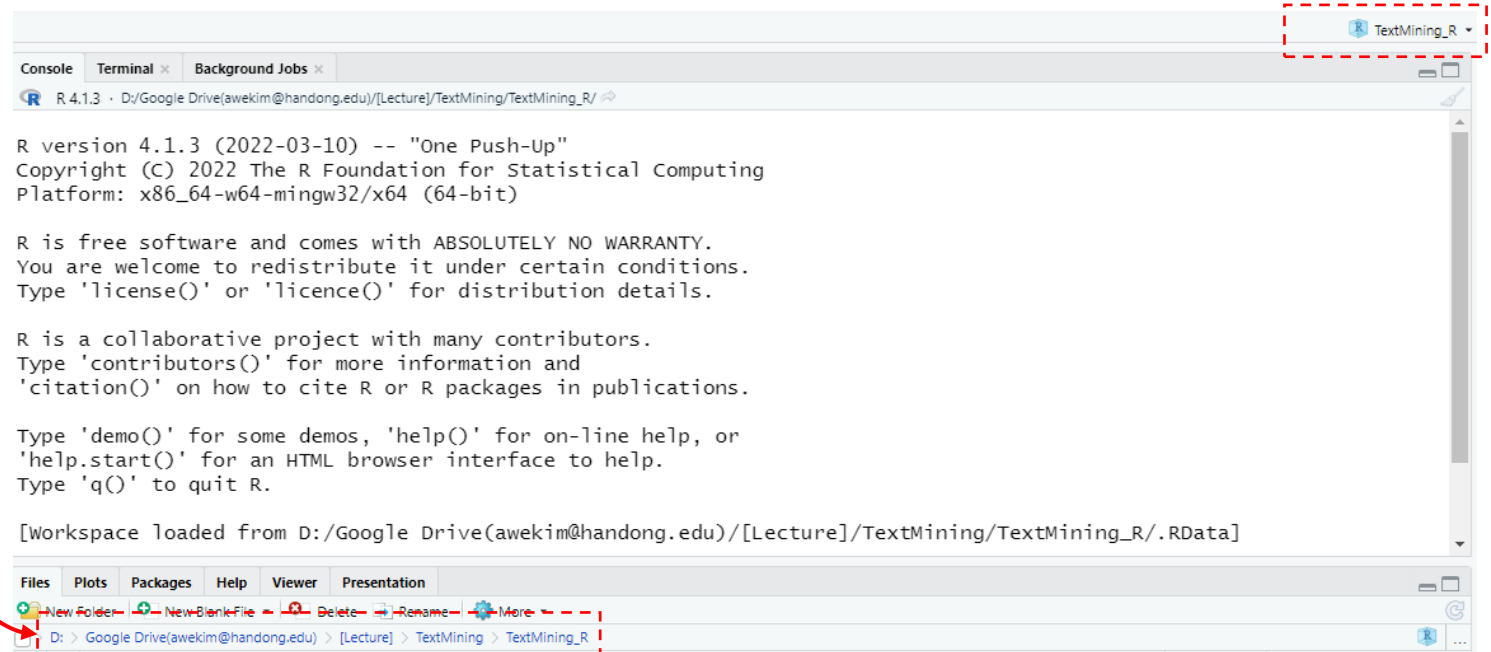

G:/내 드라이브/[Lecture]/TextMining

TextMining_R



- R project

- For each R project, the working directory is “fixed” to the folder that you created.
- With this, we can use "relative path“ easily to import data set to the R environment.



.RData file

- In R working environment, any type of R variables can be saved as .Rdata file
 - .RData file: data file used in the R working environment
- Advantages of using .RData files
 - Easy to track the changes
 - Easy to continue with the task

*raw
data*

The screenshot displays the RStudio environment with the following components:

- Source Editor:** Contains R code for downloading a file, reading a table, and exporting it.
- Console:** Shows the output of the code, including the head of the 'hotdogs' table and the results of writing to CSV and TSV files.
- Environment Pane:** Lists the variables in the workspace. The 'hotdogs' variable is highlighted, indicating it is the current object.

The 'hotdogs' variable is highlighted in the Environment pane, and a red dashed box and arrow point to it, labeled 'raw data'.

R Working Environment

- When to use .RData file?
 - Share your work in R with your colleagues
 - Continue to work on the same data on a different PC
 - Restore your work from some unexpected errors, etc.

Save .RData file – Data

- save() function
 - Used for managing variables.
 - Powerful when managing large data.

```
save(Variables, file='FileDirectory/FileName.RData')
```

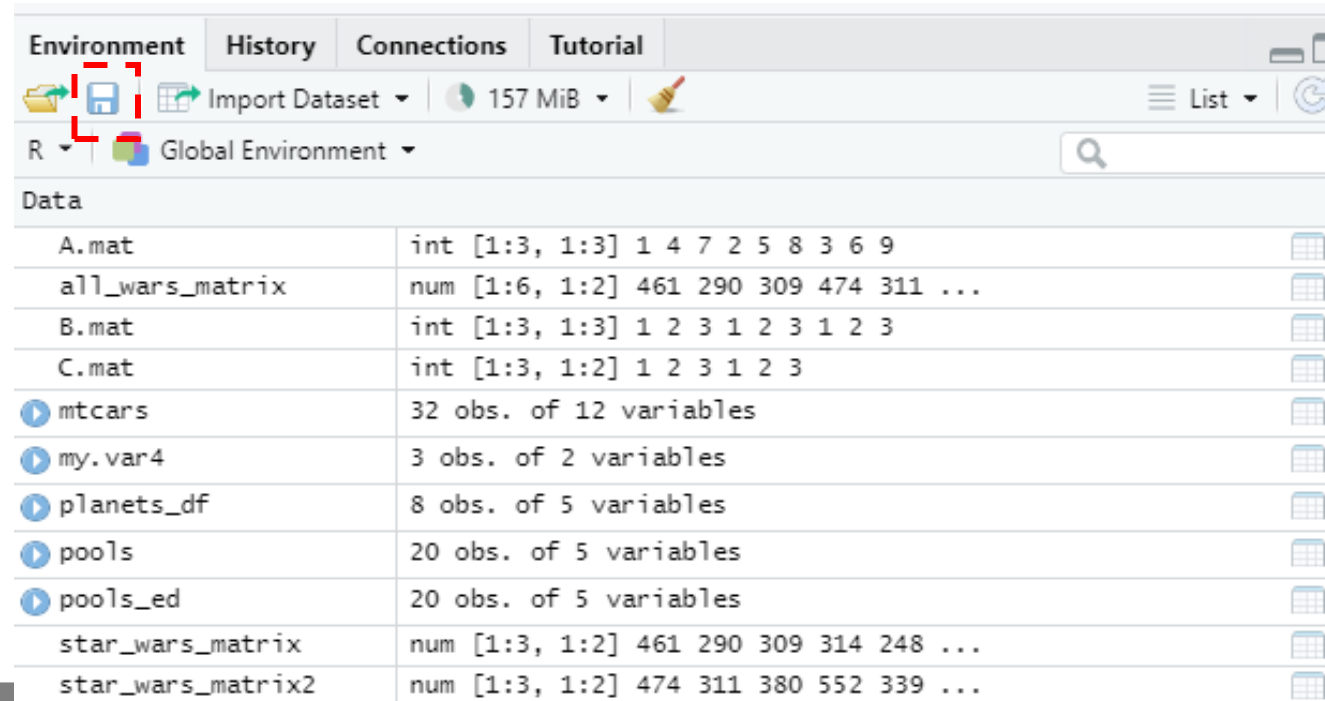
```
> my.var <- 10
> my.var2 <- c(1,4,6,22,3)
> my.var3 <- c('John', 'Bob', 'Alice')
> my.var4 <- data.frame(A = 1:3, B = 9:11)
> save(my.var, my.var2, my.var3, my.var4,
+       file = "Your directory/myVariables.RData")
```

Save .RData file – Working Environment

- Method 1) Save .Rdata of R working environment
 - Use save() function

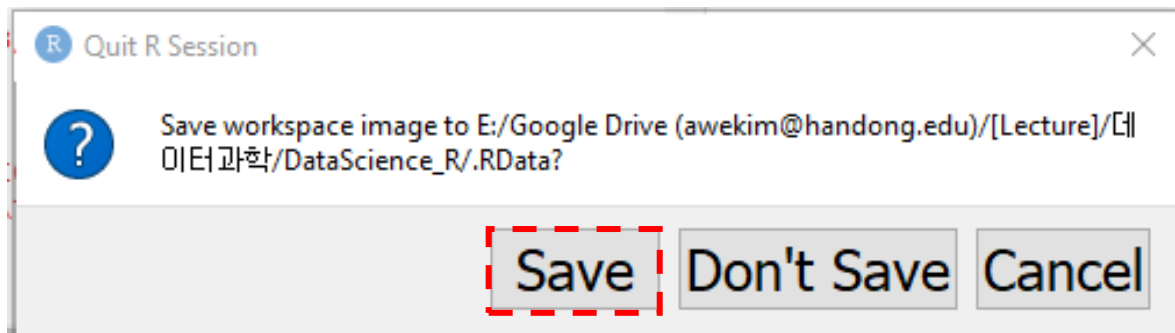
```
save(list = ls(), file='FileDirectory/FileName.RData')
```

- Press disk button



Save .RData file – Working Environment

- Method 2) Save .Rdata when closing R Session
 - By doing so, all of the tasks and variables in the R project can be saved.
 - Used for managing the whole R project



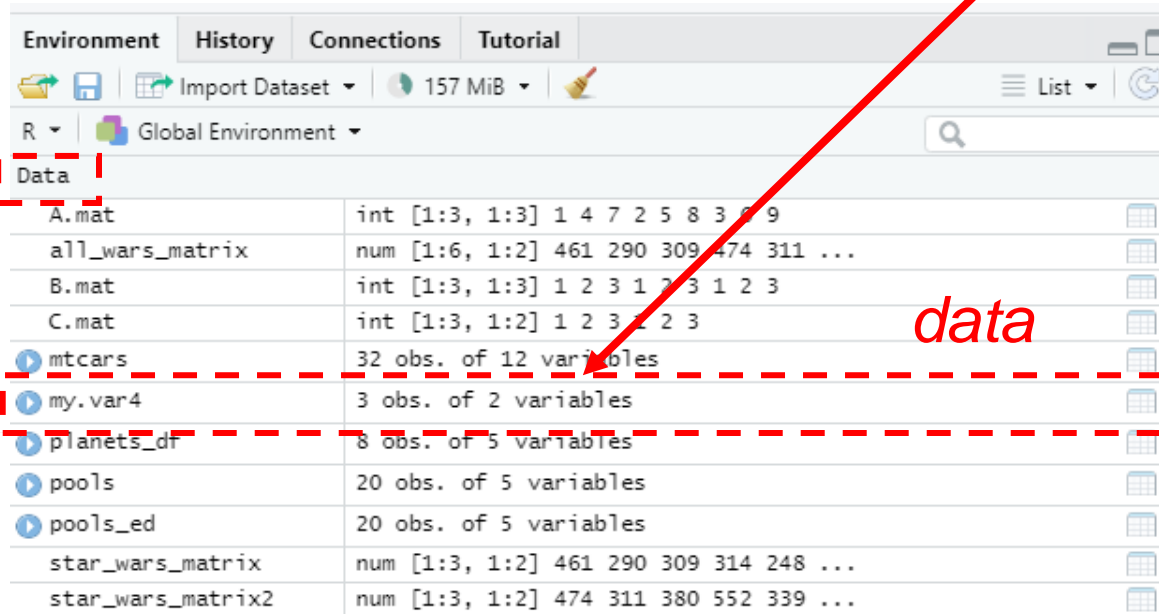
Google Drive (awekim@handong.edu) > [Lecture] > 데이터과학 > DataScience_R		
<input type="checkbox"/> Name	Type	Size
.Rproj.user	File folder	
R file	File folder	
.RData	R Workspace	
.Rhistory	RHISTORY File	
DataScience_R	R Project	
DS_check	R File	
DS_LEC02_RBASICI	R File	
DS_LEC02_RBASICII	R File	

Load .RData file

- load() function

```
load(file='FileDirectory/FileName.RData')
```

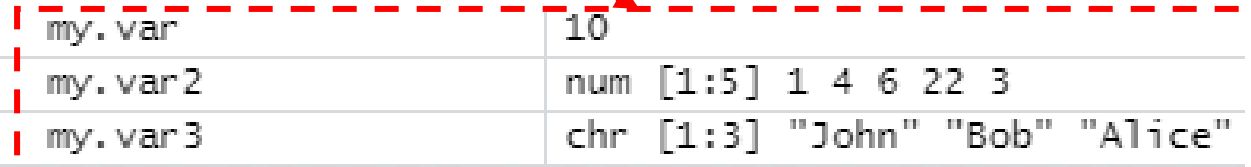
```
> load(file = "Your directory/myVariables.RData")
```



Environment	History	Connections	Tutorial
R Global Environment			
Data			
A.mat	int [1:3, 1:3]	1 4 7 2 5 8 3 6 9	
all_wars_matrix	num [1:6, 1:2]	461 290 309 474 311 ...	
B.mat	int [1:3, 1:3]	1 2 3 1 2 3 1 2 3	
C.mat	int [1:3, 1:2]	1 2 3 1 2 3	
mtcars	32 obs. of 12 variables		
my.var4	3 obs. of 2 variables		
planets_df	8 obs. of 5 variables		
pools	20 obs. of 5 variables		
pools_ed	20 obs. of 5 variables		
star_wars_matrix	num [1:3, 1:2]	461 290 309 314 248 ...	
star_wars_matrix2	num [1:3, 1:2]	474 311 380 552 339 ...	

data

value



my.var	10
my.var2	num [1:5] 1 4 6 22 3
my.var3	chr [1:3] "John" "Bob" "Alice"

Data Manipulation with dplyr I

- dplyr package
 - a grammar of data manipulation, providing a consistent set of verbs that helps you solve the most common data manipulation challenges
 - select(), filter(), mutate(), arrange(), group_by(), summarise()

```
> load(file="R file/R file_LEC02/ds_salaries_ed.RData")
> ds_sal %>% head
Error in ds_sal %>% head : could not find function "%>%"
> library(dplyr)
```

Attaching package: 'dplyr'

The following objects are masked from 'package:stats':

filter, lag

The following objects are masked from 'package:base':

intersect, setdiff, setequal, union

```
> ds_sal %>% head
```

	ID	work_year	experience_level	employment_type	job_title	salary
1	ID_0	2020	MI	FT	Data Scientist	70000
2	ID_1	2020	SE	FT	Machine Learning Scientist	260000
3	ID_2	2020	SE	FT	Big Data Engineer	85000
4	ID_3	2020	MI	FT	Product Data Analyst	20000
5	ID_4	2020	SE	FT	Machine Learning Engineer	150000
6	ID_5	2020	EN	FT	Data Analyst	72000

	salary_currency	salary_in_usd	employee_residence	remote_ratio	company_location
1	EUR	79833	DE	0	DE
2	USD	260000	JP	0	JP
3	GBP	109024	GB	50	GB
4	USD	20000	HN	0	HN
5	USD	150000	US	50	US
6	USD	72000	US	100	US

	company_size
1	L
2	S
3	M
4	S
5	L
6	L

- select()

- picks columns (variables) either with names or index
- Alternatives: (columns) indexing + slicing, etc.

*Select columns of
job_title, salary, and
salary_currency*

```
> head(ds_sal[,c("job_title", "salary", "salary_currency")])
```

```
      job_title salary salary_currency
1      Data Scientist   70000          EUR
2 Machine Learning Scientist 260000         USD
3      Big Data Engineer   85000          GBP
4      Product Data Analyst  20000          USD
5 Machine Learning Engineer 150000          USD
6      Data Analyst    72000          USD
```

*Select columns of
job_title, salary, and
salary_currency*

```
>
> ds_sal %>% select(job_title, salary, salary_currency) %>%
+   head
```

```
      job_title salary salary_currency
1      Data Scientist   70000          EUR
2 Machine Learning Scientist 260000         USD
3      Big Data Engineer   85000          GBP
4      Product Data Analyst  20000          USD
5 Machine Learning Engineer 150000          USD
6      Data Analyst    72000          USD
```

*Select columns from
job_title to
salary_currency*

```
>
> ds_sal %>% select(job_title:salary_currency) %>%
+   head
```

```
      job_title salary salary_currency
1      Data Scientist   70000          EUR
2 Machine Learning Scientist 260000         USD
3      Big Data Engineer   85000          GBP
4      Product Data Analyst  20000          USD
5 Machine Learning Engineer 150000          USD
6      Data Analyst    72000          USD
```

- select()

- picks columns (variables) either with names or index
- Alternatives: (columns) indexing + slicing, etc.

Select column number
5 and 7

```
> head(ds_sal[,c(5,7)])
      job_title salary_currency
1      Data Scientist          EUR
2 Machine Learning Scientist    USD
3      Big Data Engineer        GBP
4  Product Data Analyst         USD
5 Machine Learning Engineer    USD
6      Data Analyst            USD
```

Select column number
5 and 7

```
> ds_sal %>% select(5,7) %>%
+   head
      job_title salary_currency
1      Data Scientist          EUR
2 Machine Learning Scientist    USD
3      Big Data Engineer        GBP
4  Product Data Analyst         USD
5 Machine Learning Engineer    USD
6      Data Analyst            USD
```

Select column number
from 5 to 7

```
> ds_sal %>% select(5:7) %>%
+   head
      job_title salary salary_currency
1      Data Scientist  70000          EUR
2 Machine Learning Scientist 260000    USD
3      Big Data Engineer  85000        GBP
4  Product Data Analyst  20000         USD
5 Machine Learning Engineer 150000    USD
6      Data Analyst   72000         USD
```


dplyr::select()

- `select()` + `_with`
 - Select columns by recognizing a certain “pattern”
 - `starts_with` & `ends_with`

Select columns that start with ‘salary’

```
> ds_sal %>% select(starts_with('salary')) %>%  
+ head
```

	salary	salary_currency	salary_in_usd
1	70000	EUR	79833
2	260000	USD	260000
3	85000	GBP	109024
4	20000	USD	20000
5	150000	USD	150000
6	72000	USD	72000

Select columns that does not start with ‘salary’

```
> ds_sal %>% select(!starts_with('salary')) %>%  
+ head
```

	ID	work_year	experience_level	employment_type	job_title
1	ID_0	2020	MI	FT	Data Scientist
2	ID_1	2020	SE	FT	Machine Learning Scientist
3	ID_2	2020	SE	FT	Big Data Engineer
4	ID_3	2020	MI	FT	Product Data Analyst
5	ID_4	2020	SE	FT	Machine Learning Engineer
6	ID_5	2020	EN	FT	Data Analyst

	employee_residence	remote_ratio	company_location	company_size
1	DE	0	DE	L
2	JP	0	JP	S
3	GB	50	GB	M
4	HN	0	HN	S
5	US	50	US	L
6	US	100	US	L

Select columns that starts with “salary” or “company”?

- `select_if()`
 - Select columns by recognizing a certain “pattern”
 - Useful when selecting specific types of columns

Select columns with numeric values

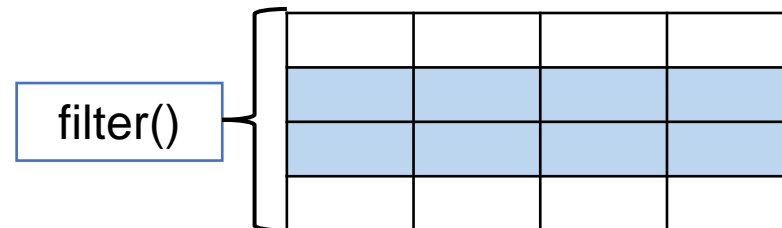
```
> ds_sal %>% select_if(is.numeric) %>% head(2)
  work_year salary salary_in_usd remote_ratio
1      2020  70000          79833           0
2      2020 260000          260000           0
```

Select columns with character values

```
> ds_sal %>% select_if(is.character) %>% head(2)
  ID experience_level employment_type job_title
1 ID_0              MI             FT Data Scientist
2 ID_1              SE             FT Machine Learning Scientist
  salary_currency employee_residence company_location company_size
1             EUR              DE              DE              L
2             USD              JP              JP              S
```

- filter()

- Subset a data frame, retaining all rows that satisfy your conditions. To be retained, the row must produce a value of 'TRUE' for all conditions.
- Alternatives: (rows) indexing + slicing, etc.



Select rows that job title is 'Data Scientist'

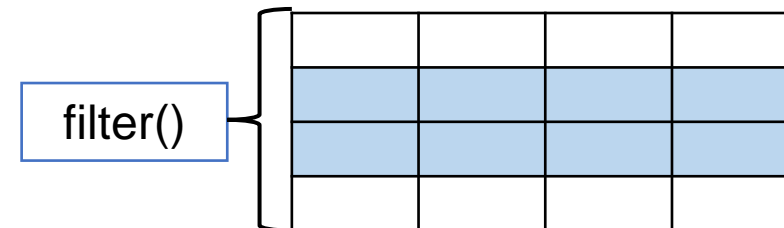
```
> head(ds_sal[ds_sal$job_title=="Data Scientist",], 2)
  ID work_year experience_level employment_type job_title salary
1 ID_0      2020             MI             FT Data Scientist  70000
8 ID_7      2020             MI             FT Data Scientist 11000000
  salary_currency salary_in_usd employee_residence remote_ratio
1              EUR       79833              DE              0
8              HUF       35735              HU              50
  company_location company_size
1              DE              L
8              HU              L
```

Select rows that job title is 'Data Scientist'

```
> ds_sal %>% filter(job_title=="Data Scientist") %>%
+   head(2)
  ID work_year experience_level employment_type job_title salary
1 ID_0      2020             MI             FT Data Scientist  70000
2 ID_7      2020             MI             FT Data Scientist 11000000
  salary_currency salary_in_usd employee_residence remote_ratio
1              EUR       79833              DE              0
2              HUF       35735              HU              50
  company_location company_size
1              DE              L
2              HU              L
```

- filter()

- Subset a data frame, retaining all rows that satisfy your conditions. To be retained, the row must produce a value of 'TRUE' for all conditions.
- Alternatives: (rows) indexing + slicing, etc.



Select rows that salary is above the average

```
> head(ds_sal[ds_sal$salary>=mean(ds_sal$salary)], 2)
  ID work_year experience_level employment_type job_title
8  ID_7      2020              MI              FT Data Scientist
12 ID_11      2020              MI              FT Data Scientist
  salary salary_currency salary_in_usd employee_residence remote_ratio
8  11000000          HUF          35735          HU          50
12 3000000          INR          40481          IN          0
  company_location company_size
8              HU              L
12             IN              L
```

Select rows that salary is above the average

```
> ds_sal %>% filter(salary>=mean(salary)) %>%
+   head(2)
  ID work_year experience_level employment_type job_title
1  ID_7      2020              MI              FT Data Scientist
2 ID_11      2020              MI              FT Data Scientist
  salary salary_currency salary_in_usd employee_residence remote_ratio
1 11000000          HUF          35735          HU          50
2 3000000          INR          40481          IN          0
  company_location company_size
1              HU              L
2             IN              L
```

- arrange()
 - Orders the rows of a data frame by the values of selected columns
 - arrange(desc(x)) : arrange in a descending order with x
 - Alternatives: order

Arrange the data set in an ascending order with salary

```
> head(ds_sal[order(ds_sal$salary),],2)
      ID work_year experience_level employment_type job_title
186 ID_185      2021              MI              FT Data Engineer
239 ID_238      2021              EN              FT Data Scientist
      salary salary_currency salary_in_usd employee_residence remote_ratio
186   4000          USD          4000          IR          100
239   4000          USD          4000          VN           0
      company_location company_size
186              IR              M
239              VN              M
```

Arrange the data set in an ascending order with salary

```
> ds_sal %>% arrange(salary) %>% head(2)
      ID work_year experience_level employment_type job_title salary
1 ID_185      2021              MI              FT Data Engineer  4000
2 ID_238      2021              EN              FT Data Scientist  4000
      salary_currency salary_in_usd employee_residence remote_ratio
1          USD          4000          IR          100
2          USD          4000          VN           0
      company_location company_size
1              IR              M
2              VN              M
```

- `arrange()`
 - Orders the rows of a data frame by the values of selected columns
 - `arrange(desc(x))` : arrange in a descending order with `x`
 - Alternatives: `order`

Arrange the data set in a descending order with salary

```
> head(ds_sal[order(ds_sal$salary, decreasing=TRUE),],2)
      ID work_year experience_level employment_type job_title
178 ID_177      2021              MI             FT Data Scientist
8   ID_7        2020              MI             FT Data Scientist
      salary salary_currency salary_in_usd employee_residence
178 30400000             CLP          40038                CL
8   11000000             HUF          35735                HU
      remote_ratio company_location company_size
178           100             CL             L
8           50             HU             L
```

Arrange the data set in a descending order with salary

```
> ds_sal %>% arrange(desc(salary)) %>% head(2)
      ID work_year experience_level employment_type job_title
1 ID_177      2021              MI             FT Data Scientist
2 ID_7        2020              MI             FT Data Scientist
      salary salary_currency salary_in_usd employee_residence remote_ratio
1 30400000             CLP          40038                CL          100
2 11000000             HUF          35735                HU           50
      company_location company_size
1             CL             L
2             HU             L
```

Data Manipulation with dplyr II

- mutate()
 - **Adds new variables** that are functions of existing variables
 - Use ifelse (if & else), case_when (if & elseif & else) with mutate()

ifelse(condition, TRUE-result, FALSE-result)

<Convert numeric variable to dummy variable>

Create salary.d variable, which returns High if salary_in_usd is higher than its average, and Low otherwise

Create experience variable referencing year 2022.

```
> ds_sal %>% mutate(experience=2023-work_year) %>%  
+   select(work_year,experience,salary) %>% head
```

	work_year	experience	salary
1	2020	3	70000
2	2020	3	260000
3	2020	3	85000
4	2020	3	20000
5	2020	3	150000
6	2020	3	72000

```
> ds_sal %>% mutate(salary.d =  
+   ifelse(salary_in_usd > mean(salary_in_usd),  
+   "High", "Low")) %>%  
+   select(work_year,salary,salary.d) %>% head
```

	work_year	salary	salary.d
1	2020	70000	Low
2	2020	260000	High
3	2020	85000	Low

- mutate()
 - **Adds new variables** that are functions of existing variables
 - Use ifelse (if & else), case_when (if & elseif & else) with mutate()

Create International variable, which returns International if employee_residence and company_location are the different and Domestic otherwise

```
> ds_sal %>% mutate(International =  
+   ifelse(employee_residence != company_location,  
+         "International","Domestic")) %>%  
+   select(employee_residence,company_location,International) %>%  
+   head
```

	employee_residence	company_location	International
1	DE	DE	Domestic
2	JP	JP	Domestic
3	GB	GB	Domestic
4	HN	HN	Domestic
5	US	US	Domestic
6	US	US	Domestic

- mutate()
 - **Adds new variables** that are functions of existing variables
 - Use ifelse (if & else), case_when (if & elseif & else) with mutate()

Create job.d variable, which returns DS for Data Scientist, DA for Data Analyst and others for Others

```
> ds_sal %>%  
+   mutate(job.d = case_when(job_title=="Data Scientist" ~ "DS",  
+                             job_title=="Data Analyst" ~ "DA",  
+                             TRUE ~ "Others")) %>%  
+   select(work_year, job_title, job.d) %>% head  
  work_year      job_title  job.d  
1      2020      Data Scientist    DS  
2      2020 Machine Learning Scientist Others  
3      2020      Big Data Engineer Others  
4      2020      Product Data Analyst Others  
5      2020 Machine Learning Engineer Others  
6      2020      Data Analyst    DA
```

- mutate_at()
 - changes selected columns

<Log Transformation>

Convert all variables into a log scale

```
> ds_sal %>% select(ID, salary, experience) %>%  
+   mutate_at(vars(salary, experience), log) %>% head
```

	ID	salary	experience
1	ID_0	11.156251	1.098612
2	ID_1	12.468437	1.098612
3	ID_2	11.350407	1.098612
4	ID_3	9.903488	1.098612
5	ID_4	11.918391	1.098612
6	ID_5	11.184421	1.098612

- mutate_all()
 - changes all columns

```
> ds_sal %>%
+   mutate_all(is.na) %>% head(2)
  ID work_year experience_level employment_type job_title salary
1 FALSE      FALSE           FALSE           FALSE    FALSE  FALSE
2 FALSE      FALSE           FALSE           FALSE    FALSE  FALSE
  salary_currency salary_in_usd employee_residence remote_ratio
1                FALSE        FALSE              FALSE        FALSE
2                FALSE        FALSE              FALSE        FALSE
  company_location company_size experience salary.d job.d
1                FALSE        FALSE        FALSE  FALSE FALSE
2                FALSE        FALSE        FALSE  FALSE FALSE
```

<Normalization>

Normalize numeric variables

Z-score Normalization

$$\frac{x - \text{mean}(x)}{\text{std.dev}(x)}$$

```
> norm.fun <-
+   function(x){
+     (x - mean(x, na.rm = TRUE)) / sd(x, na.rm = TRUE)}
> ds_sal %>% select_if(is.numeric) %>%
+   mutate_all(norm.fun) %>% head
  work_year      salary salary_in_usd remote_ratio
1 -2.030349 -0.16446973  -0.45752711  -1.7421785
2 -2.030349 -0.04144122   2.08156475  -1.7421785
3 -2.030349 -0.15475696  -0.04613862  -0.5139528
4 -2.030349 -0.19684566  -1.30075303  -1.7421785
5 -2.030349 -0.11266825   0.53133577  -0.5139528
6 -2.030349 -0.16317470  -0.56791751   0.7142729
```

dplyr::rename()

- `rename()`
 - `rename()` changes the names of individual variables using *new_name = old_name*

```
> names(ds_sal)
[1] "ID" "work_year" "experience_level"
[4] "employment_type" "job_title" "salary"
[7] "salary_currency" "salary_in_usd" "employee_residence"
[10] "remote_ratio" "company_location" "company_size"
[13] "experience" "salary.d" "job.d"
> library(magrittr)
> # rename
> names(ds_sal)
[1] "ID" "work_year" "experience_level"
[4] "employment_type" "job_title" "salary"
[7] "salary_currency" "salary_in_usd" "employee_residence"
[10] "remote_ratio" "company_location" "company_size"
[13] "experience" "salary.d" "job.d"
> ds_sal %<>% rename(sal.type=salary.d,
+                  job.type=job.d)
> names(ds_sal)
[1] "ID" "work_year" "experience_level"
[4] "employment_type" "job_title" "salary"
[7] "salary_currency" "salary_in_usd" "employee_residence"
[10] "remote_ratio" "company_location" "company_size"
[13] "experience" "sal.type" "job.type"
```

- `rename_with()`
 - rename columns with functions

```
> ds_sal %>% rename_with(toupper) %>% names
[1] "ID" "WORK_YEAR" "EXPERIENCE_LEVEL"
[4] "EMPLOYMENT_TYPE" "JOB_TITLE" "SALARY"
[7] "SALARY_CURRENCY" "SALARY_IN_USD" "EMPLOYEE_RESIDENCE"
[10] "REMOTE_RATIO" "COMPANY_LOCATION" "COMPANY_SIZE"
[13] "EXPERIENCE" "SAL.TYPE" "JOB.TYPE"
> ds_sal %>% rename_with(toupper, starts_with("salary")) %>% names
[1] "ID" "work_year" "experience_level"
[4] "employment_type" "job_title" "SALARY"
[7] "SALARY_CURRENCY" "SALARY_IN_USD" "employee_residence"
[10] "remote_ratio" "company_location" "company_size"
[13] "experience" "sal.type" "job.type"
```

- group_by()

- group_by() takes an existing tbl and converts it into a grouped tbl where operations are performed "by group".
- ungroup() removes grouping

```
> ds_sal_gr <- ds_sal %>% group_by(job_title)
> ds_sal_gr
# A tibble: 607 x 15
# Groups:   job_title [50]
   ID work_year experience_level employment_type job_title salary
  <chr>   <int>   <chr>           <chr>      <chr>    <int>
1 ID_0    2020 MI             FT      Data Scientist 7 e4
2 ID_1    2020 SE             FT      Machine Learni~ 2.6 e5
3 ID_2    2020 SE             FT      Big Data Engin~ 8.5 e4
4 ID_3    2020 MI             FT      Product Data A~ 2 e4
5 ID_4    2020 SE             FT      Machine Learni~ 1.5 e5
6 ID_5    2020 EN             FT      Data Analyst 7.2 e4
7 ID_6    2020 SE             FT      Lead Data Scie~ 1.9 e5
8 ID_7    2020 MI             FT      Data Scientist 1.1 e7
9 ID_8    2020 MI             FT      Business Data ~ 1.35e5
10 ID_9   2020 SE             FT      Lead Data Engi~ 1.25e5
# ... with 597 more rows, and 9 more variables: salary_currency <chr>,
# salary_in_usd <int>, employee_residence <chr>, remote_ratio <int>,
# company_location <chr>, company_size <chr>, experience <dbl>,
# salary.d <chr>, job.d <chr>
```

```
> ds_sal_gr %>% ungroup
# A tibble: 607 x 15
   ID work_year experience_level employment_type job_title salary
  <chr>   <int>   <chr>           <chr>      <chr>    <int>
1 ID_0    2020 MI             FT      Data Scientist 7 e4
2 ID_1    2020 SE             FT      Machine Learni~ 2.6 e5
3 ID_2    2020 SE             FT      Big Data Engin~ 8.5 e4
4 ID_3    2020 MI             FT      Product Data A~ 2 e4
5 ID_4    2020 SE             FT      Machine Learni~ 1.5 e5
6 ID_5    2020 EN             FT      Data Analyst 7.2 e4
7 ID_6    2020 SE             FT      Lead Data Scie~ 1.9 e5
8 ID_7    2020 MI             FT      Data Scientist 1.1 e7
9 ID_8    2020 MI             FT      Business Data ~ 1.35e5
10 ID_9   2020 SE             FT      Lead Data Engi~ 1.25e5
# ... with 597 more rows, and 9 more variables: salary_currency <chr>,
# salary_in_usd <int>, employee_residence <chr>, remote_ratio <int>,
# company_location <chr>, company_size <chr>, experience <dbl>,
# salary.d <chr>, job.d <chr>
```

dplyr::group_by() + mutate()

- `group_by() + mutate()`
 - summarise the data based on the “group_by” target

```
> ds_sal %>%  
+ # group_by(company_size) %>%  
+   mutate(salary.mean=mean(salary)) %>%  
+   data.frame %>% head(2)
```

	ID	work_year	experience_level	employment_type	job_title	salary	salary_currency	salary_in_usd
1	ID_0	2020	MI	FT	Data Scientist	70000	EUR	79833
2	ID_1	2020	SE	FT	Machine Learning Scientist	260000	USD	260000

	employee_residence	remote_ratio	company_location	company_size	experience	salary.d	job.d	salary.mean
1	DE	0	DE	L	3	Low	DS	324000.1
2	JP	0	JP	S	3	High	Others	324000.1

average of all samples

```
> ds_sal %>%  
+ group_by(company_size) %>%  
+   mutate(salary.mean=mean(salary)) %>%  
+   data.frame %>% head(2)
```

	ID	work_year	experience_level	employment_type	job_title	salary	salary_currency	salary_in_usd
1	ID_0	2020	MI	FT	Data Scientist	70000	EUR	79833
2	ID_1	2020	SE	FT	Machine Learning Scientist	260000	USD	260000

	employee_residence	remote_ratio	company_location	company_size	experience	salary.d	job.d	salary.mean
1	DE	0	DE	L	3	Low	DS	593695.8
2	JP	0	JP	S	3	High	Others	377710.0

average of each group

dplyr::group_by() + summarise()

- group_by() + summarise()
 - summarise the data based on the “group_by” target

```
> ds_sal_gr %>% summarise(salary=mean(salary))
# A tibble: 50 x 2
  job_title          salary
  <chr>             <dbl>
1 3D Computer Vision Researcher 400000
2 AI Scientist                290571.
3 Analytics Engineer          175000
4 Applied Data Scientist      172400
5 Applied Machine Learning Scientist 141350
6 BI Data Analyst            1902045.
7 Big Data Architect          125000
8 Big Data Engineer           455000
9 Business Data Analyst       355000
10 Cloud Data Engineer        140000
# ... with 40 more rows
> ds_sal_gr %>% ungroup %>%
+ summarise(salary=mean(salary))
# A tibble: 1 x 1
  salary
  <dbl>
1 324000.
```

```
> ds_sal %>%
+ group_by(job_title, company_size) %>%
+ dplyr::summarise(salary=mean(salary))
`summarise()` has grouped output by 'job_title'. You can override
`.groups` argument.
# A tibble: 98 x 3
# Groups:   job_title [50]
  job_title          company_size salary
  <chr>             <chr>     <dbl>
1 3D Computer Vision Researcher M        400000
2 AI Scientist       L        127500
3 AI Scientist       M        66000
4 AI Scientist       S        549000
5 Analytics Engineer M        175000
6 Applied Data Scientist L        172400
7 Applied Machine Learning Scientist L        249000
8 Applied Machine Learning Scientist M         33700
9 BI Data Analyst    L       5575000
10 BI Data Analyst    M         99000
# ... with 88 more rows
```


Data Visualization with ggplot

- ggplot2() package
 - A well-known package for creating data visualizations
 - Recommended to take “Data Visualization” to learn more details

`data %>% ggplot(aes(x= xvar, y=yvar)) +`
`geom_XX() + XX`

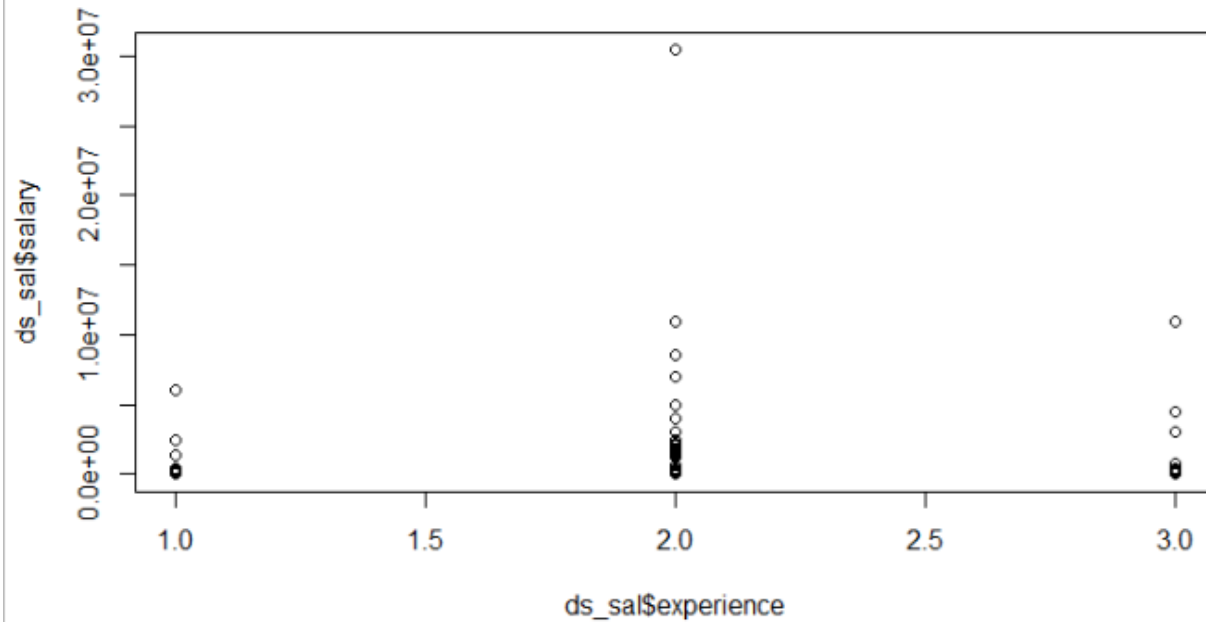
Define x and y variables

Define visualization type

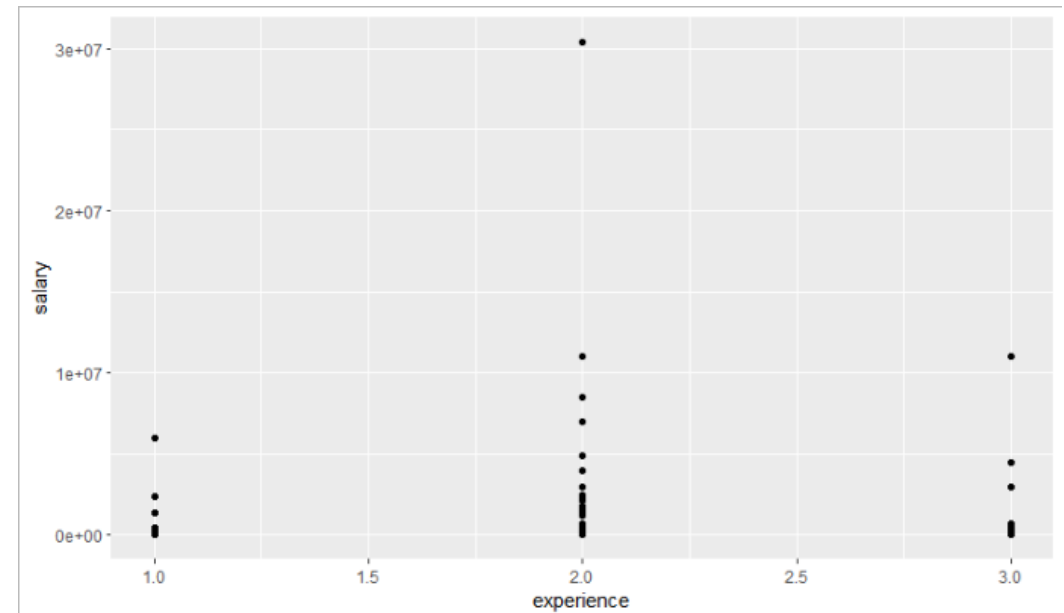
Add more options if necessary

- Scatter plot
 - `geom_point()`

```
> plot(ds_sal$experience, ds_sal$salary)
```



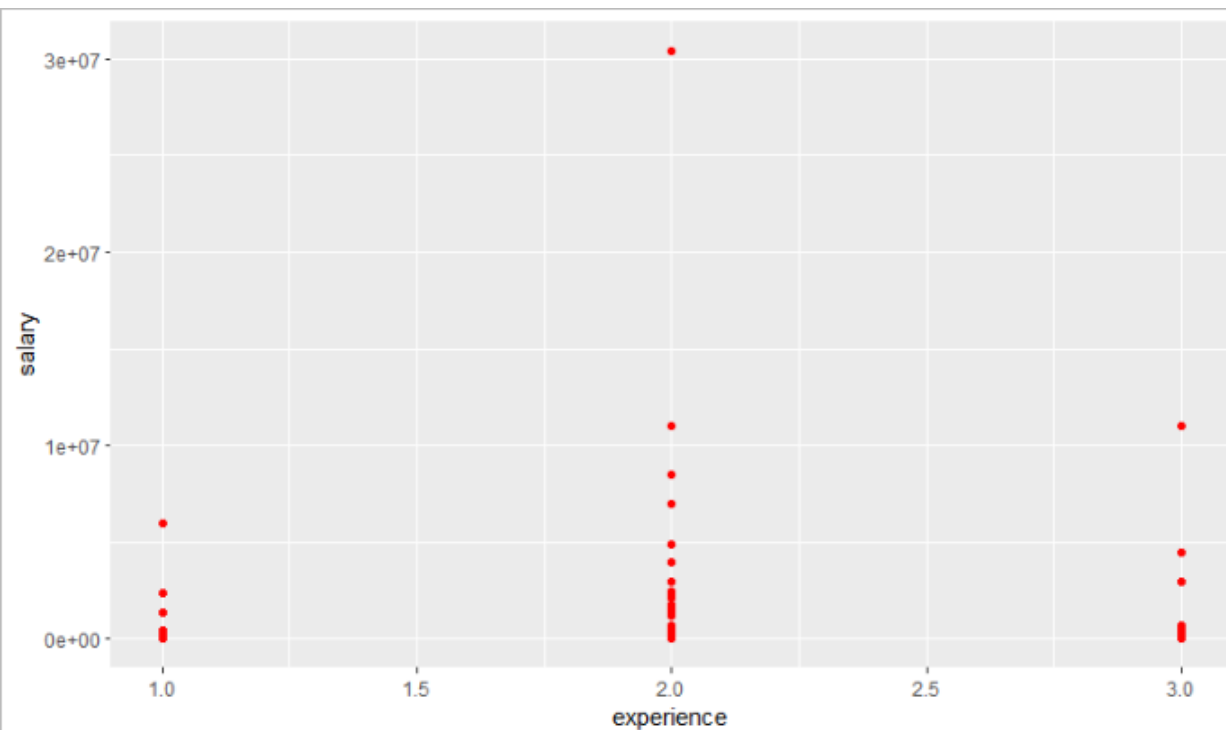
```
> ds_sal %>% ggplot(aes(x=experience, y=salary)) +  
  + geom_point()
```



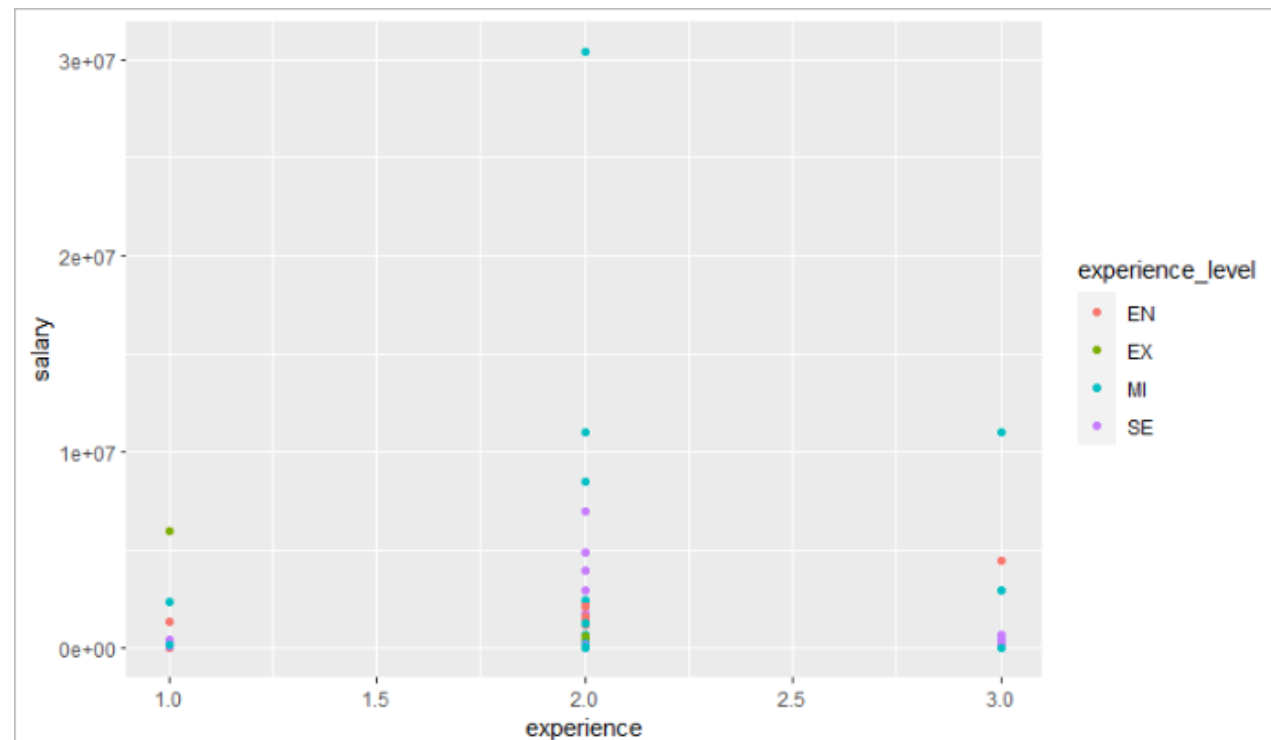
- Scatter plot

- `geom_point()` + `color`

```
> ds_sal %>%  
+   ggplot(aes(x=experience, y=salary)) +  
+   geom_point(color="red")
```




```
> ds_sal %>%  
+   ggplot(aes(x=experience, y=salary,  
+             color=experience_level)) +  
+   geom_point()
```



- Line plot


- geom_line()

```
> ds_sal %>% group_by(work_year) %>%  
+ summarise(salary_in_usd = mean(salary_in_usd)) %>%  
+ ggplot(aes(x=work_year, y=salary_in_usd)) +  
+ geom_line()
```



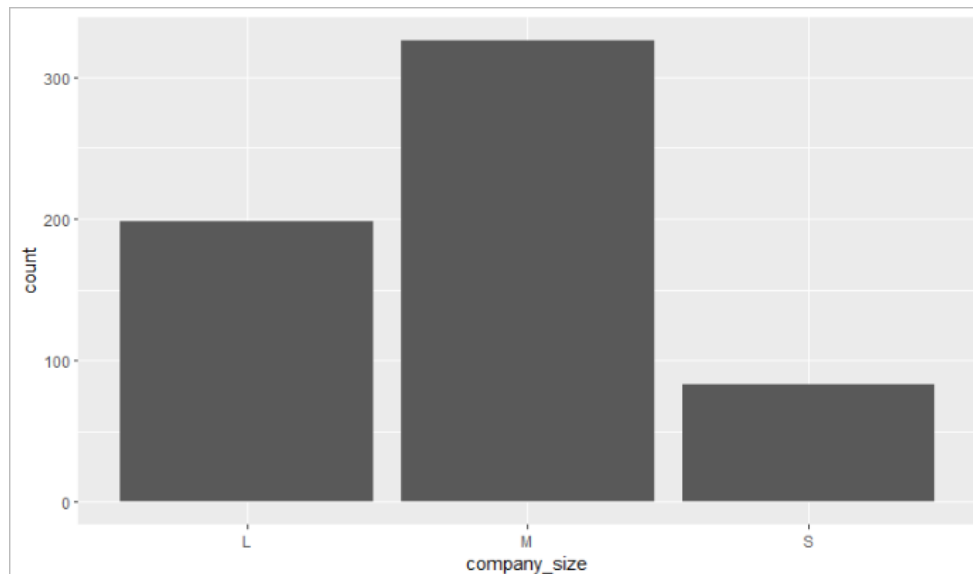
EN Entry-level / Junior MI Mid-level /
Intermediate SE Senior-level / Expert EX Executive-level

```
> ds_sal %>% group_by(work_year, experience_level) %>%  
+ summarise(salary_in_usd = mean(salary_in_usd)) %>%  
+ ggplot(aes(x=work_year, y=salary_in_usd,  
+ group=experience_level,  
+ color=experience_level)) +  
+ geom_line()
```

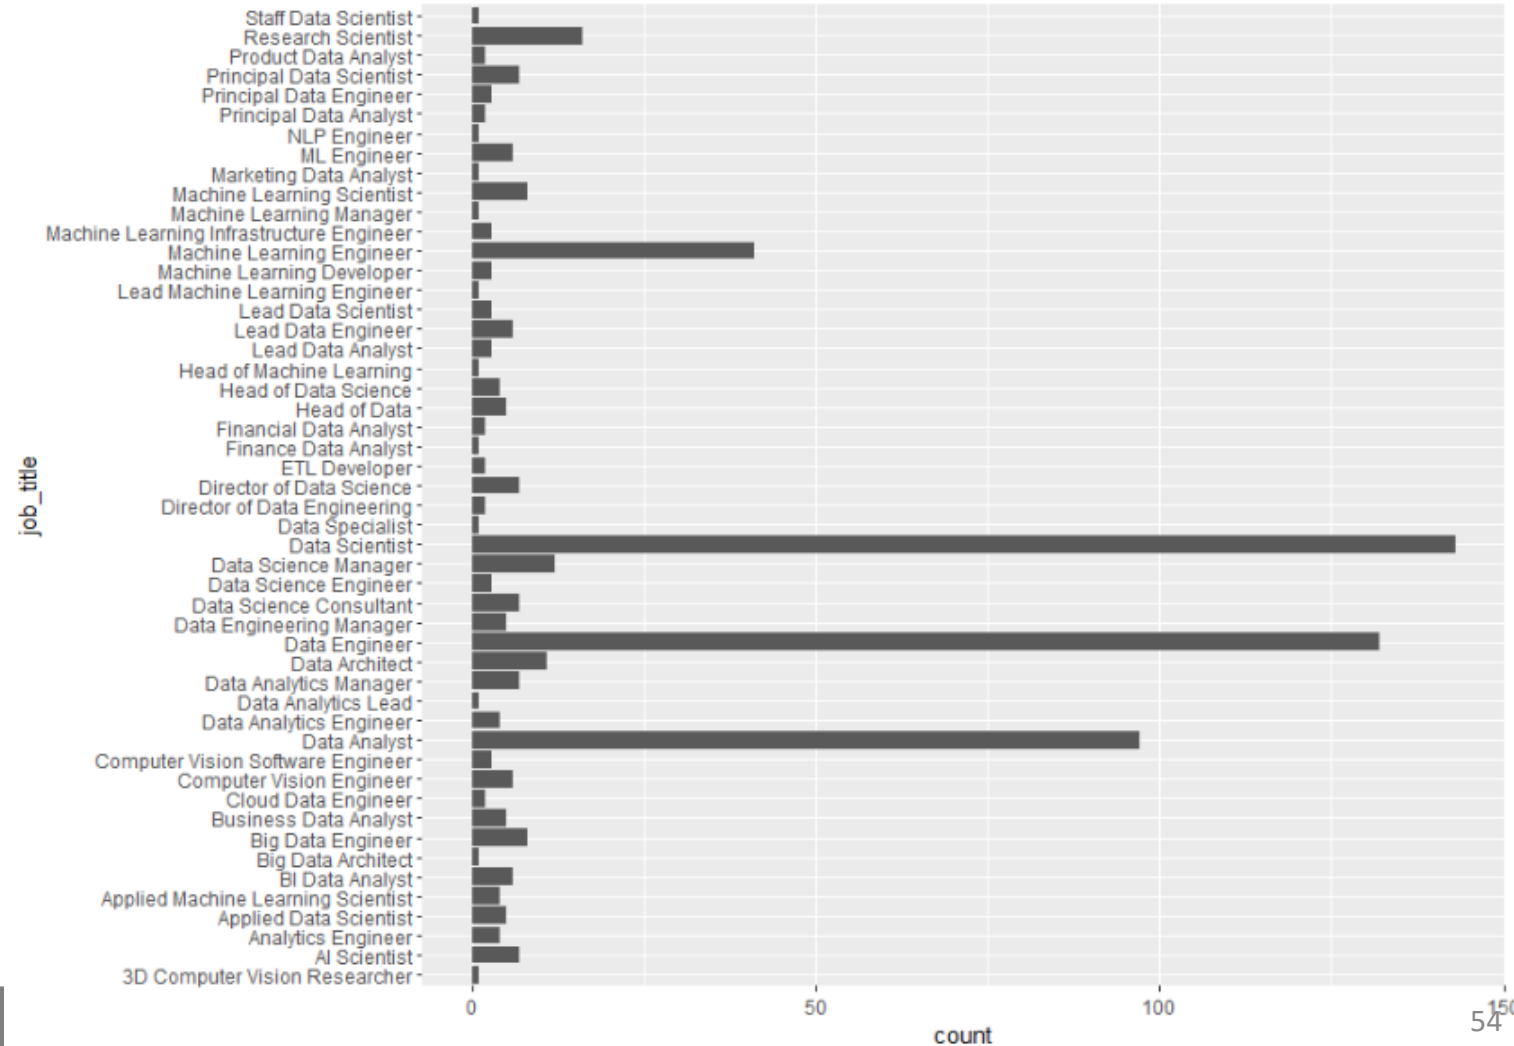


- Bar plot - frequency
 - geom_bar()

```
> ds_sal %>% ggplot(aes(company_size)) +  
+   geom_bar()
```

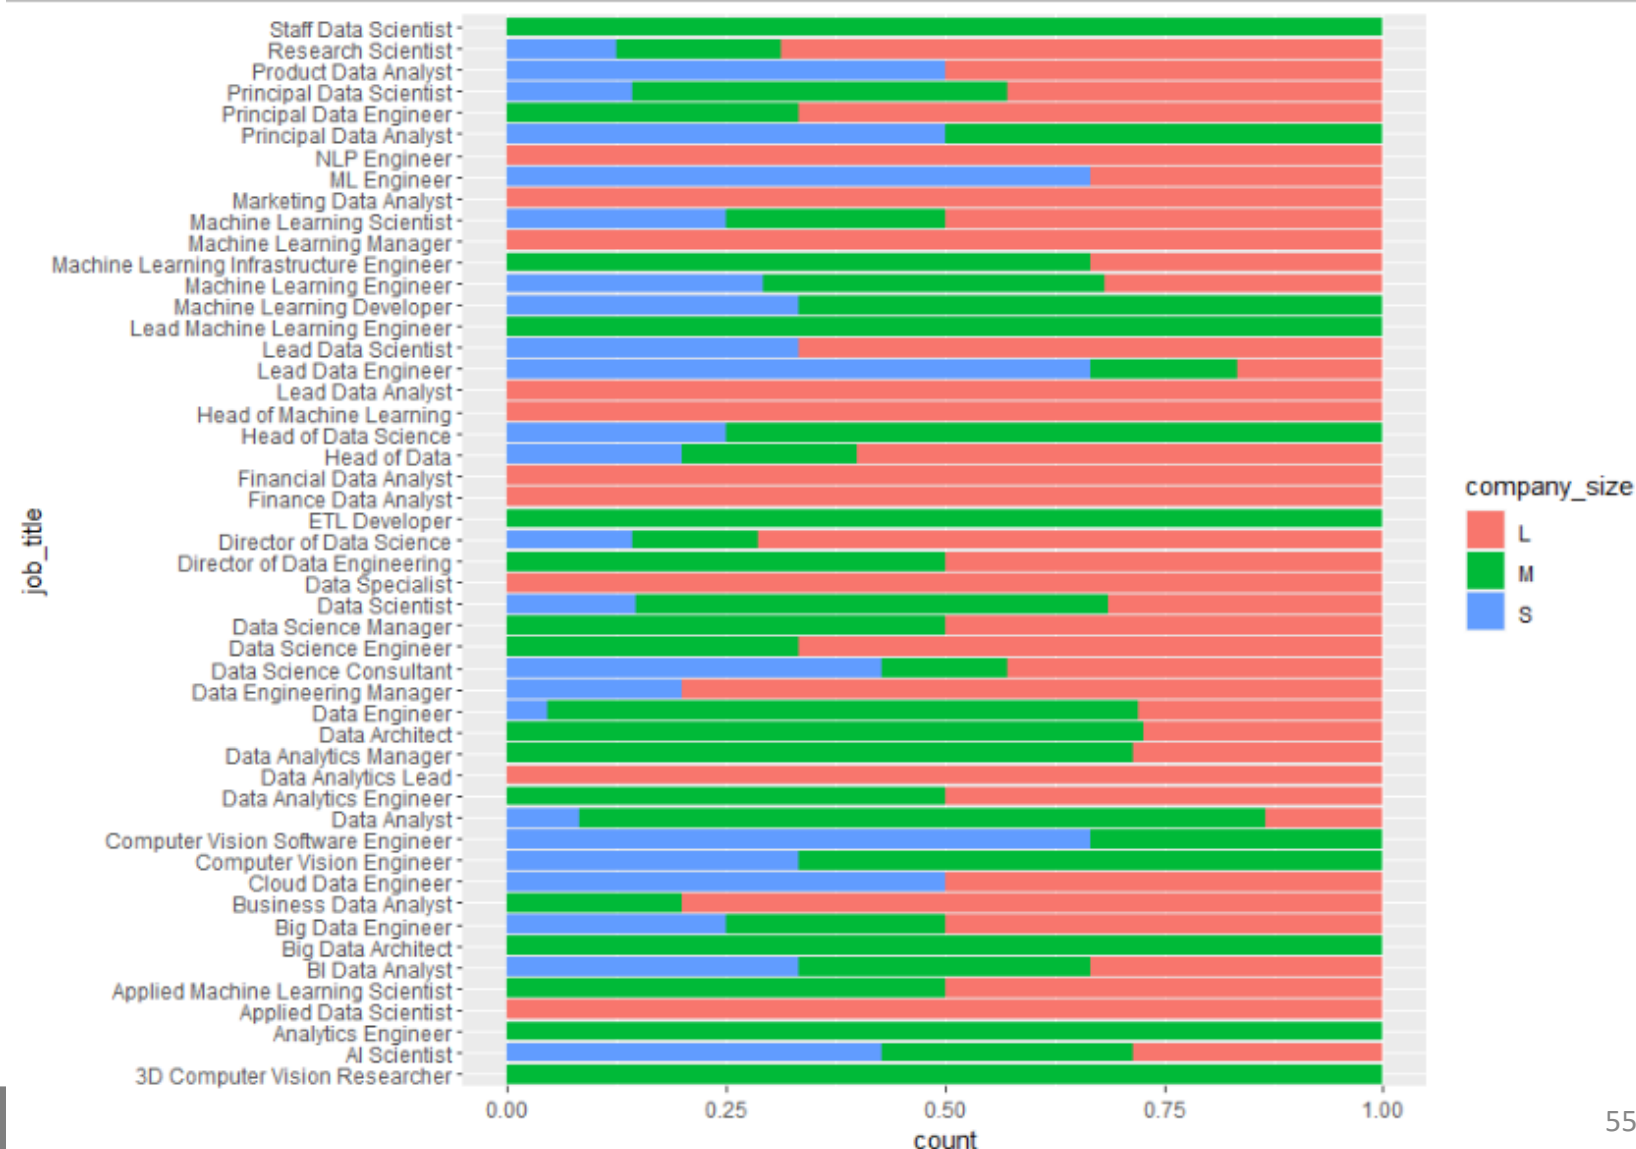


```
> ds_sal %>% ggplot(aes(job_title)) +  
+   geom_bar() + coord_flip()
```




- Bar plot - frequency per group
 - `geom_bar()`

```
> ds_sal %>%
+   ggplot(aes(x=job_title, fill=company_size)) +
+   geom_bar(position='fill') + coord_flip()
```




- Bar plot - frequency
 - geom_bar()

```
> ds_sal %>% group_by(job.d, company_size) %>%  
+   summarise(salary=mean(salary)) %>%  
+   ggplot(aes(x=job.d, y=salary, fill=company_size)) +  
+   geom_bar(stat='identity')
```



```
> ds_sal %>% group_by(job.d, company_size) %>%  
+   summarise(salary=mean(salary)) %>%  
+   ggplot(aes(x=job.d, y=salary, fill=company_size)) +  
+   geom_bar(stat='identity', position='dodge')
```



Possible Errors

- Version issue

```
> library(ggplot2)
```

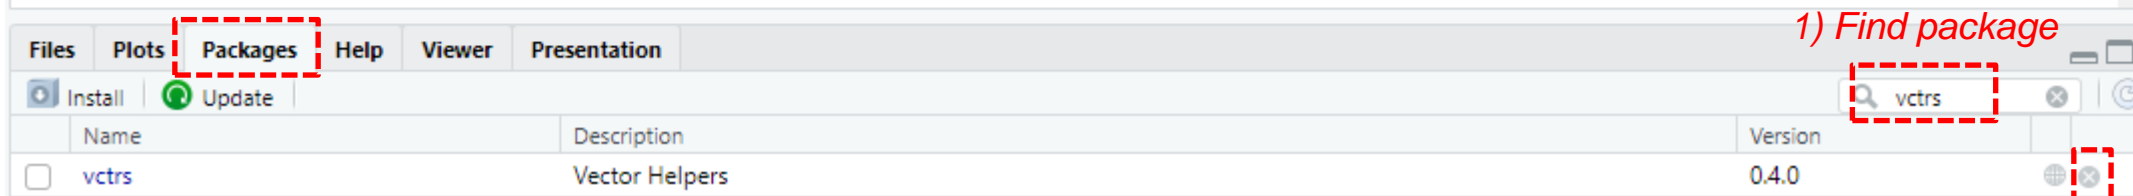
```
Error: package or namespace load failed for 'ggplot2' in loadNamespace  
(i, c(lib.loc, .libPaths()), versionCheck = vI[[i]]):
```

네임스페이스 'vctrs' 0.4.0는 이미 로드되었으나 $\geq 0.5.0$ 가 필요합니다

```
In addition: Warning message:
```

```
패키지 'ggplot2'는 R 버전 4.1.3에서 작성되었습니다
```

```
> |
```



3) `install.packages('vctrs')`